

# eda

June 2, 2025

```
[55]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go

%matplotlib inline
```

```
[56]: data_path = r"data\DigitalExposome Dataset.csv"
df = pd.read_csv(data_path)
```

```
[57]: df.columns
```

```
[57]: Index(['IBI', 'HR', 'NO2', 'Noise', 'NH3', 'PM10', 'CO', 'PM25', 'Label',
          'PM1', 'EDA', 'BVP'],
          dtype='object')
```

```
[58]: df.head()
```

```
[58]:
```

	IBI	HR	NO2	Noise	NH3	PM10	CO	PM25	\
0	0.0	0.377574	0.0	0.511358	0.003018	0.003091	0.871758	0.000000	
1	0.0	0.196398	0.0	0.490903	0.003018	0.003091	0.876848	0.003091	
2	0.0	0.454163	0.0	0.470449	0.006036	0.006181	0.881939	0.006181	
3	0.0	0.322451	0.0	0.449995	0.009055	0.009272	0.887030	0.009272	
4	0.0	0.237595	0.0	0.429540	0.012073	0.012362	0.892121	0.012362	

	Label	PM1	EDA	BVP
0	5	0.000000	0.0	0.0
1	5	0.001854	0.0	0.0
2	5	0.003709	0.0	0.0
3	5	0.005563	0.0	0.0
4	5	0.007417	0.0	0.0

```
[59]: df.shape
```

```
[59]: (42436, 12)
```

```
[60]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42436 entries, 0 to 42435
Data columns (total 12 columns):
#   Column  Non-Null Count  Dtype
---  -
0   IBI      42436 non-null   float64
1   HR       42436 non-null   float64
2   NO2      42436 non-null   float64
3   Noise    42436 non-null   float64
4   NH3      42436 non-null   float64
5   PM10     42436 non-null   float64
6   CO       42436 non-null   float64
7   PM25     42436 non-null   float64
8   Label    42436 non-null   int64
9   PM1      42436 non-null   float64
10  EDA      42436 non-null   float64
11  BVP      42436 non-null   float64
dtypes: float64(11), int64(1)
memory usage: 3.9 MB
```

```
[61]: df.describe().T
```

```
[61]:
```

	count	mean	std	min	25%	50%	75%	max
IBI	42436.0	0.178688	0.230127	0.0	0.000000	0.016760	0.309984	1.0
HR	42436.0	0.530687	0.266531	0.0	0.309505	0.491213	0.776388	1.0
NO2	42436.0	0.502239	0.274914	0.0	0.285412	0.489703	0.736983	1.0
Noise	42436.0	0.427541	0.247601	0.0	0.214008	0.403671	0.618592	1.0
NH3	42436.0	0.564276	0.294515	0.0	0.305285	0.644756	0.826943	1.0
PM10	42436.0	0.208474	0.210981	0.0	0.017624	0.158454	0.326889	1.0
CO	42436.0	0.679874	0.308111	0.0	0.458047	0.805352	0.909832	5.0
PM25	42436.0	0.224901	0.223242	0.0	0.012362	0.176150	0.376813	1.0
Label	42436.0	3.326185	1.599863	1.0	2.000000	4.000000	5.000000	5.0
PM1	42436.0	0.230773	0.255253	0.0	0.006310	0.176316	0.407268	24.0
EDA	42436.0	0.260318	0.221725	0.0	0.060686	0.216064	0.421109	1.0
BVP	42436.0	0.482902	0.175569	0.0	0.449963	0.512280	0.579756	1.0

```
[62]: missing = df.isna().sum()
missing[missing > 0].sort_values(ascending=False)
```

```
[62]: Series([], dtype: int64)
```

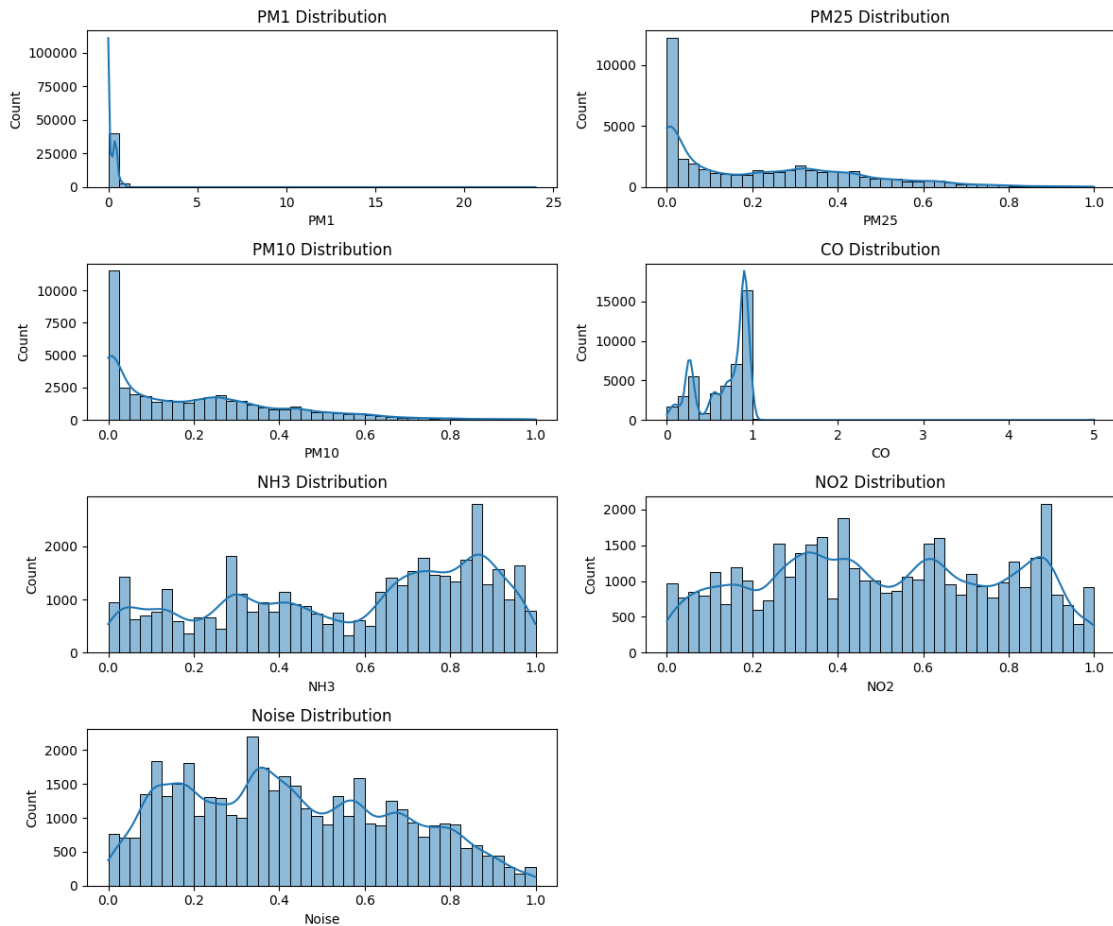
```
[ ]: env_cols = ["PM1", "PM25", "PM10", "CO", "NH3", "NO2", "Noise"]

plt.figure(figsize=(10, 8))
for i, col in enumerate(env_cols, 1):
    plt.subplot(4, 2, i)
```

```

sns.histplot(df[col], kde=True, bins=40)
plt.title(f"{col} Distribution")
plt.tight_layout()
plt.show()

```



```
[64]: (df["PM1"] > 1.0).sum(), (df["PM1"] > 5.0).sum())
```

```
[64]: (np.int64(1), np.int64(1))
```

```

[65]: sample = df.sample(5000, random_state=0)
fig = px.density_heatmap(sample,
                        x="PM25",
                        y="PM10",
                        nbinsx=40,
                        nbinsy=40,
                        title="Hexbin: PM25 vs. PM10 (normalized)",
                        width=800,
                        height=600

```

```
)
fig.show()
```

```
[66]: def noise_category(x):
        if x < 0.3:
            return "Low"
        elif x < 0.6:
            return "Medium"
        else:
            return "High"

df["Noise_cat"] = df["Noise"].apply(noise_category)
fig = px.box(
    df,
    x="Noise_cat",
    y="EDA",
    title="EDA by Noise Category",
    width=800,
    height=600
)
fig.show()
```

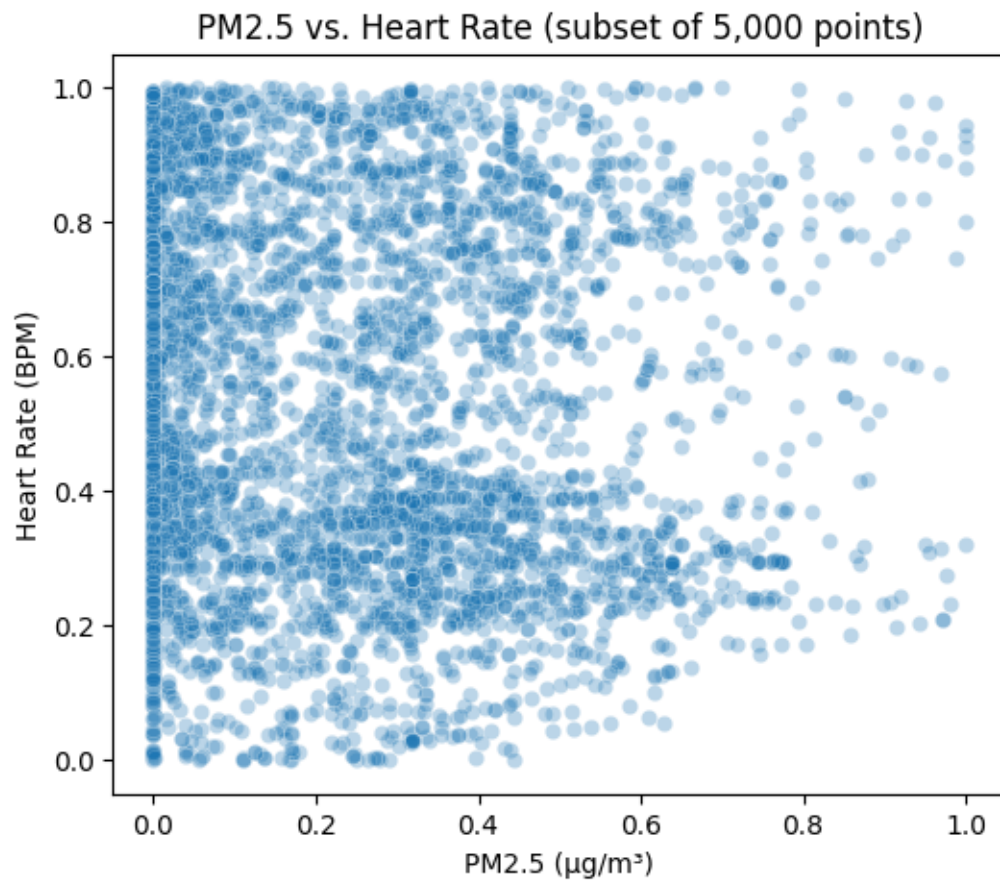
```
[67]: corr = ␣
        ↪df[["PM1", "PM25", "PM10", "CO", "NH3", "NO2", "Noise", "HR", "EDA", "BVP", "IBI"]].
        ↪corr()
fig = px.imshow(
    corr,
    text_auto=".2f",
    title="Correlation Matrix (All Env + Physio Variables)",
    width=800,
    height=600
)
fig.show()
```

```
[68]: spikes = df[df["PM1"] > 5.0]
fig = go.Figure()
fig.add_trace(go.Scatter(x=spikes.index, y=spikes["PM1"], mode="markers", ␣
    ↪name="PM1"))
fig.add_trace(go.Scatter(x=spikes.index, y=spikes["HR"], mode="lines", name="HR",
    ↪(closeup)))
fig.update_layout(
    title="PM1 Spikes & HR Overlaps",
    xaxis_title="Time",
    yaxis_title="Values",
    width=800,
    height=600
)
```


```
fig.show()
```

```
[69]: df["PM25_bin"] = pd.qcut(df["PM25"], q=4, labels=False)
fig = px.box(
    df,
    x="PM25_bin",
    y="EDA",
    title="EDA by PM2.5 Quartile",
    width=800,
    height=600
)
fig.show()
```

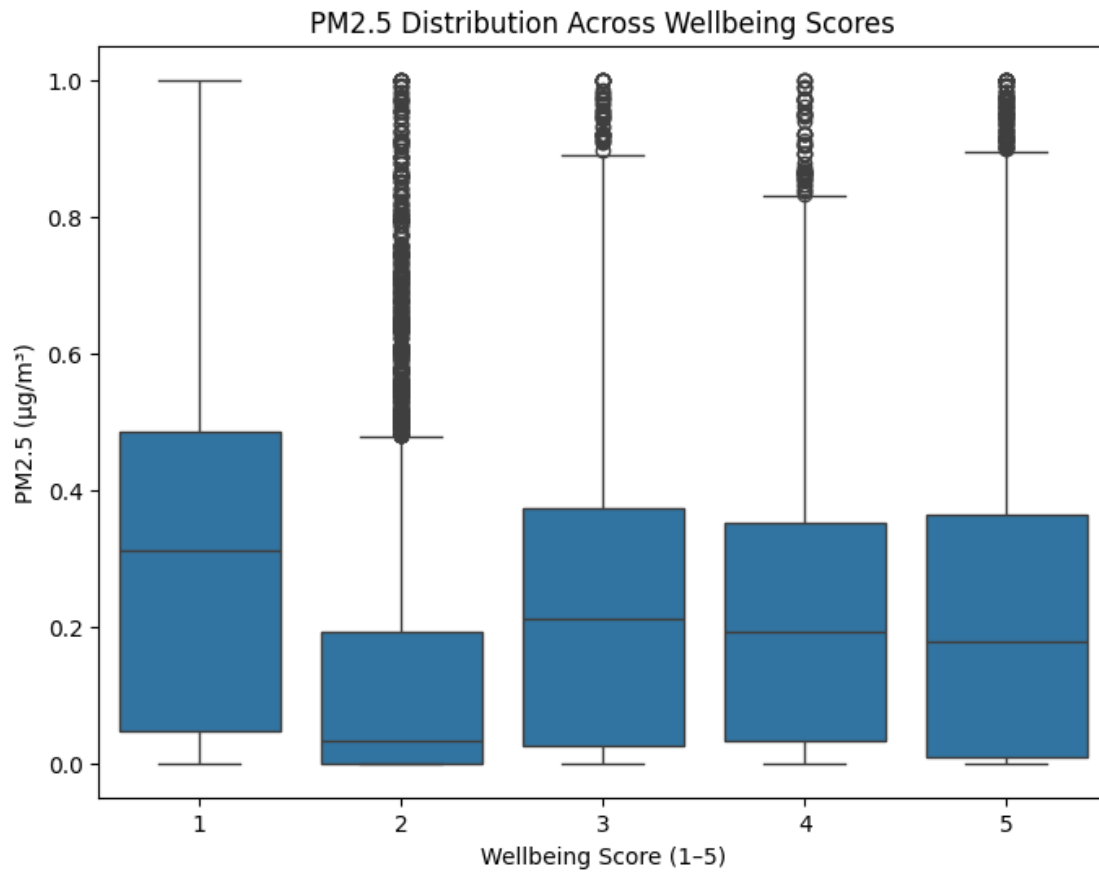
```
[70]: plt.figure(figsize=(6, 5))
sns.scatterplot(x="PM25", y="HR", data=df.sample(5000), alpha=0.3)
plt.title("PM2.5 vs. Heart Rate (subset of 5,000 points)")
plt.xlabel("PM2.5 ( $\mu\text{g}/\text{m}^3$ )")
plt.ylabel("Heart Rate (BPM)")
plt.show()
```



```
[71]: sns.jointplot(x="PM25", y="HR", data=df, kind="hex", gridsize=30)
plt.suptitle("Hexbin: PM2.5 vs. HR", y=1.02)
plt.show()
```



```
[72]: plt.figure(figsize=(8, 6))
sns.boxplot(x="Label", y="PM25", data=df)
plt.title("PM2.5 Distribution Across Wellbeing Scores")
plt.xlabel("Wellbeing Score (1-5)")
plt.ylabel("PM2.5 (µg/m³)")
plt.show()
```



[ ]: